

Fish Group — AI Architecture & Deployment Options

A guide to the different ways Fish Group can interact with Claude and the codebase, from local development to fully autonomous cloud agents.

AWS EC2 · Claude Code · OpenClaw · March 2026 · Prepared for Michael Fisch

The Tool Spectrum: From Workflow to Autonomous

Before choosing a deployment model, it helps to understand where each tool sits on the agency spectrum.

Tool	Who defines steps?	Runs 24/7?	Best for	~Cost/mo
n8n	You — visually	Yes (server)	Deterministic triggers	\$80–150
Claude Code	You set goal; AI plans	No (needs human)	Builds, code gen, dev work	\$20–200
OpenClaw	Agent autonomously	Yes (cloud)	Autonomous repeatable ops	\$100–200

The winning strategy: use all three. Claude Code builds and maintains the skills. OpenClaw runs them autonomously. n8n handles webhook triggers if needed.

Claude Seat Pricing — 2 Seats (Michael + Emil)

Both Michael and Emil need Claude Code and Cowork access. All paid plans include both features — the difference is token limits and reset windows. For active Claude Code development, Pro's 5-hour window caps out fast during coding sessions.

Plan	Price/seat	Token Limit	Reset	Claude Code	Cowork	Best for
Pro	\$20/mo	~44K tokens / ~45 msgs	Every 5 hrs	Yes	Yes	Business analysts
Max 5x	\$100/mo	~88K tokens (5x Pro)	Weekly	Yes	Yes	Developers & leads
Max 20x	\$200/mo	~220K tokens (20x Pro)	Weekly	Yes	Yes	Power users

Recommendation for Fish Group: Start with **1x Max 5x + 1x Pro = \$120/mo total**. Only one person is likely deep in Claude Code at a time — share the dev load so the Max seat goes to whoever is actively building. The Pro seat covers the other for writing, analysis, and lighter Cowork tasks. Upgrade the second seat to Max 5x if both are coding simultaneously.

Deployment Options

Option 1 — Local Dev (Claude Code on Laptop)

GOOD START

Michael or Emil runs Claude Code locally from their MacBook / PC

The simplest entry point. Install Claude Code on your machine, give it access to your APIs (Airtable, Google, etc.) and work with it like a very smart pair programmer. You describe what you want, it builds it, you review.

PROS	CONS
+ Zero server cost	- Not always-on — laptop must be open
+ Instant setup (you already have it)	- No autonomous execution
+ Full control — you're in the loop on every step	- API keys stored locally
+ Best for exploration and building	- Doesn't scale to client-facing automation

Cost: \$20/mo Claude Max · \$0 server · Best for: building skills & one-time tasks

Option 2 — Claude Professional with Seats

LIMITED

Multiple Claude.ai Pro accounts for Michael, Emil, and team

Good for writing, summarizing, and conversational AI — but **not suitable for automated workflows or agents**. No API access, no MCP servers, no persistent agents.

PROS	CONS
+ Easy to provision team seats	- No API access — can't connect to Airtable/Google/etc.
+ No technical setup required	- No persistent agents or cron jobs
+ Good for writing, analysis, Q&A;	- Not suitable for automation
+ Clients can get seats too	- \$20-25/seat/month adds up quickly

Cost: \$20–25/seat/mo · Best for: team writing & analysis — not automation

Option 3 — EC2 + OpenClaw (Always-On Agent)

RECOMMENDED

Fish Group internal agent running 24/7 on AWS EC2

OpenClaw runs continuously on EC2, listening for webhook triggers and executing cron jobs without anyone needing to be at a laptop. This is the Finn agent system we scoped in the session.

```
Fish Group AWS Architecture (Option 3) ----- [ Michael /
Emil - Claude Code CLI ] | v [ EC2 t3.small ~$17/mo ] OpenClaw Gateway (:18789) |-- Finn
(orchestrator) |-- Client Ops agent |-- Data/Airtable agent |-- Permissions agent AWS KMS /
OpenRouter | [ Google Workspace / Airtable AWS IAM / QuickBooks / ShipStation ]
```

PROS	CONS
+ Always-on — runs cron jobs, listens for webhooks	- Requires initial server setup (~2hr)
+ Secrets stored securely in AWS KMS	- Small ongoing EC2 cost
+ Scales to multiple clients (one instance each)	- Need SSH access to debug/update
+ OpenRouter models — cheaper than Anthropic direct	
+ One AWS account per client = isolated billing	

Cost: \$17–35/mo EC2 · \$20–50/mo OpenRouter · ~\$5/mo KMS · Total: ~\$50–90/mo per instance

Option 4 — Hybrid: Local Claude Code + EC2 OpenClaw

RECOMMENDED

The full Fish Group architecture — build locally, run autonomously in cloud

Michael and Emil use Claude Code on their laptops to **build and edit skills**. Those skills live in a Git repo. EC2 runs OpenClaw autonomously. This is Greg's personal setup — best of both worlds.

```
Hybrid Architecture ----- [ Michael + Emil — Claude Code
(local) ] git push -> GitHub repo | git pull v [ EC2 Fish Group ] [ EC2 Per Client ] OpenClaw +
Finn OpenClaw onboarding, Piermont / Gary's permissions, briefs Local dev + EC2 = fast builds +
always-on execution.
```

PROS	
+ Local dev = fast iteration, no server round trips	
+ EC2 = always-on execution, no laptop dependency	
+ Git as source of truth — easy rollback	CONS
+ Emil can dev locally while agent runs on EC2	- Slightly more setup than a single option
+ Client EC2 accounts isolated — their bill, their data	- Requires Git workflow discipline

Cost: \$20/mo Claude Code · \$50–90/mo EC2 Fish Group · \$25–50/mo per client instance

Our Recommendation for Fish Group

Start with Option 1 to get hands dirty. Move to Option 4 within 2–4 weeks.

Phase 1 — Now (this week)

- Michael + Emil: use Claude Code locally — build first 2 skills (client onboarding, Airtable lookup)
- Set up one AWS account for Fish Group internal with per-client account strategy
- Generate Airtable personal access token + Google Workspace OAuth

Phase 2 — Follow-up Session (~2 weeks)

- Deploy OpenClaw on EC2 for Fish Group internal (Finn goes live)
- Connect Google Workspace, Airtable, AWS IAM via secrets
- Wire up first autonomous workflow: new client onboarding
- Set up morning ops brief cron job

Phase 3 — Piermont + Gary's (4–8 weeks)

- Piermont gets own EC2 instance with QuickBooks + ShipStation + Airtable agent
- Build chat UI portal for Piermont
- Scope Gary's CS agent — email triage first, voice AI second